

CS4025: Words

- Words
- Bayesian Reasoning with words:
 - Spelling correction
 - Sentiment Analysis

Reading: J&M (5.1 to 5.6 in 1st ed, 3.10 to 3.11 in 2nd)

Spelling correction

Fred's dog likes to chase our cag.

- What is the misspelled word?
- What should it be?

Spelling errors

- 80% of all spelling errors are
 - ♦ character insertion (*the* -> *ther*)
 - ♦ character deletion (*the* -> *th*)
 - ♦ character substitution (*the* -> *thw*)
 - ♦ character transposition (*the* -> *teh*)
- This defines a set of candidates
 - ♦ *cag* <- *car*, *cat*, *cage*, *cga*, ...

Example: *acress*

Which words in the dictionary can be obtained by a single transformation?

Error	Correction	Transformation			
		Correct Letter	Error Letter	Position (Letter #)	Type
acress	actress	t	—	2	deletion
acress	cress	—	a	0	insertion
acress	caress	ca	ac	0	transposition
acress	access	c	r	2	substitution
acress	across	o	e	3	substitution
acress	acres	—	2	5	insertion
acress	acres	—	2	4	insertion

How choose between these?

- Statistical model of error frequency
 - » *e*->*r* more common than *e* -> *p* (keyboard)
- Statistical model of words
 - » *cat* is more common than *cage*
- AI: world knowledge, what's plausible
 - » dogs like to chase cats

A combined statistical model

The noisy channel model:



What comes out is affected by what went in and how the channel distorts it.

The model precisely

- Assume we have received the noisy word O
- We seek to choose the possible input w which maximises $P(w|O)$
- $P(w|O)$ is the probability that w was intended, given that O was received
- By Bayes' rule, this is equivalent to

$$P(w|O) = \frac{P(O|w) \cdot P(w)}{P(O)}$$

The model precisely - 2

- We know O , and need to find the value of w that maximises:

$$P(w|O) = \frac{P(O|w) \cdot P(w)}{P(O)}$$

- Since $P(O)$ is the same for all w , we just need to maximise:

$$P(O|w) \cdot P(w)$$

- $P(w)$ is the prior (depends on the *language* model)
 - Predict more likely words
- $P(O|w)$ is the conditional likelihood (depends on the *channel* model)
 - Predict more likely mis-spellings

The language model

- How do we estimate $P(w)$ for a possible word w ?
- We collect a large corpus of text and see what proportion of words are w :

$$\frac{\text{Number of times } w \text{ appears}}{\text{Total number of words}}$$

- In practice, we need to “smooth” the counts to cater for the fact that a possible word may not appear in the corpus (see J&M)

The channel model

- How likely are the various transformations of the original word?

Option 1: Use a corpus of errors to estimate types of errors (Birkbeck spelling error corpus <http://ota.ahds.ac.uk/texts/0643.html>)

- Count e.g.
 - $\text{trans}[x,y]$ is how often the characters xy were typed as yx
 - $\text{count}[x,y]$ is how often xy occurs in the correct words
- If O is the result of transposing x and y :
 - $P(O|w) = \text{trans}[x,y] / \text{count}[x,y]$

The channel model

- How likely are the various transformations of the original word?

Option 2: Use your knowledge of typing and language to create the channel model.

- Assign higher probability to words which require fewer transforms
- Vowel substitution is more common than consonant substitution
- Making a mistake on the first letter is less likely
- Errors between adjacent letters on keyboard are more likely
- Insertion errors are more likely if the same character is repeated
-

Example

c	freq(c)	p(c)	p(1 c)	p(1 c)p(c)	%
actress	1343	.0000315	.000117	3.69×10^{-9}	37%
ciess	0	.000000014	.00000144	2.02×10^{-14}	0%
caress	4	.0000001	.00000164	1.64×10^{-13}	0%
access	2280	.000058	.000000209	1.21×10^{-11}	0%
across	8436	.00019	.0000093	1.77×10^{-9}	18%
acres	2879	.000065	.0000321	2.09×10^{-9}	21%
acres	2879	.000065	.0000342	2.22×10^{-9}	23%

- % is the result of dividing by P(**acress**) and multiplying by 100
- Note that **acres** can be obtained in 2 ways and so will win overall.

Algorithm 1 for single errors

function SpellingChecker (Dictionary, Text)

1) **Load** Dictionary (as hash table usually)

2) **For each word in** text

 (a) **If** word **in** Dictionary

 (i) **do** nothing

Else

 (i) **apply** noisy channel model to correct

Dealing with multiple errors

- If we assume errors are independent, if there are several we can just multiply their probabilities.
- But if the words are very different, it may be hard to align them and determine the set of errors made.
- We need to find the minimum set of operations needed to transform one word into another – the *minimum edit distance*. Then we can compute $P(\text{O}|\text{w})$ using this set.
- There are different ways to visualise such alignments:

Dealing with multiple errors

Trace

i	n	t	e	n	t	i	o	n
/	/	/	/					
e	x	e	c	u	t	i	o	n

Alignment

i	n	t	e	n	g	t	i	o	n
g	e	x	e	c	u	t	i	o	n

Operation

List

	delete i →	i	n	t	e	n	t	i	o	n
	substitute n by e →	n	t	e	n	t	i	o	n	
	substitute t by x →	e	t	e	n	t	i	o	n	
	insert u →	e	x	e	n	t	i	o	n	
	substitute n by c →	e	x	e	n	u	t	i	o	n
		e	x	e	c	u	t	i	o	n

Finding the minimal edit distance

- Assume that each possible edit has a cost, e.g. insertions and deletions cost 1, substitutions have a cost of 2. (Costs can depend on the characters).
- E.g. the alignment on the last slide scores 8. Is this the best alignment?
- Maintain an array *edit-distance* such that:
 - » $edit-distance[i,j]$ = the minimal distance (sum of costs) between the first i characters of the target and the first j characters of the source.
- Each cell can be computed as a function of the surrounding cells:

Example

Source

n	9	10	11	10	11	12	11	10	9	8
o	8	9	10	9	10	11	10	9	8	9
i	7	8	9	8	9	10	9	8	9	10
t	6	7	8	7	8	9	8	9	10	11
n	5	6	7	6	7	8	9	10	11	10
e	4	3	6	5	6	7	8	9	10	11
t	3	4	5	6	7	8	7	10	11	12
n	2	3	4	5	6	7	8	8	10	9
i	1	2	3	4	5	6	7	6	9	10
#	0	1	2	3	4	5	6	7	8	9
	#	e	x	e	c	u	t	i	o	n

Target

The idea

- An alignment is a path from $[0,0]$ to $[m,n]$ through adjacent cells. \nearrow = subst, \uparrow = del, \rightarrow = insert
- Best path through cell “?” could involve:

X	?
Y	Z

Passing through Y and then substituting
Passing through X and then inserting, or
Passing through Z and then deleting

- Best score for this cell is **$\min(Y+\text{subst}, X+\text{ins}, Z+\text{del})$**
where ins and del are 1, subst is 0 or 2, depending on whether the characters are the same or not

The algorithm (sketch)

```
function min-edit-distance(target, source) -> min-distance
    n = length(target); m = length(source);
    create a distance matrix distance[n+1, m+1];
    distance[0, 0] = 0;
    for all i, distance[i, 0] = ins-cost * i;
    for all j, distance[0, j] = del-cost * j;
    for each column i from 1 to n do
        for each row j from 1 to m do
            distance[i, j] =
                min(distance[i-1, j] + ins-cost,
                    distance[i, j-1] + del-cost,
                    distance[i-1, j-1]
                        + subst-cost(source[j], target[i]))
    return distance[n, m];
```

Example

Source

n	9	10	11	10	11	12	11	10	9	8
o	8	9	10	9	10	11	10	9	8	9
i	7	8	9	8	9	10	9	8	9	10
t	6	7	8	7	8	9	8	9	10	11
n	5	6	7	6	7	8	9	10	11	10
e	4	3	6	5	6	7	8	9	10	11
t	3	4	5	6	7	8	7	10	11	12
n	2	3	4	5	6	7	8	8	10	9
i	1	2	3	4	5	6	7	6	9	10
#	0	1	2	3	4	5	6	7	8	9
	#	e	x	e	c	u	t	i	o	n

Target

Reading off the solution

n	9	10	11	10	11	12	11	10	9	8
o	8	9	10	9	10	11	10	9	8	9
i	7	8	9	8	9	10	9	8	9	10
t	6	7	8	7	8	9	8	9	10	11
n	5	6	7	6	7	8	9	10	11	10
e	4	3	6	5	6	7	8	9	10	11
t	3	4	5	6	7	8	7	10	11	12
n	2	3	4	5	6	7	8	8	10	9
i	1	2	3	4	5	6	7	6	9	10
#	0	1	2	3	4	5	6	7	8	9
	#	e	x	e	c	u	t	i	o	n

Real-world errors

- What if a word is misspelled as another word?

- » Typographical errors

- insertion, deletion, substitution and transposition

- *buckled* for *bucked*, *his* for *this*

- » Homophones

- *dessert* for *desert*, *piece* for *peace*

*Fred's dog likes to chase our **chat***

- » Legal, but not very plausible

- more likely *chat* is misspelling of *cat*

- » How detect this? Need to assess plausibility of the sentence

- Calculate Likelihood of sequences of words (ngrams)

First Practical: Sentiment Analysis

TASK: Given a textual review of a movie, we need to decide if it is **Positive** or **Negative**.

- it's so laddish and juvenile , only teenage boys could possibly find it funny .
- take care of my cat offers a refreshingly different slice of asian cinema .
- interesting but not compelling
- everytime you think undercover brother has run out of steam , it finds a new way to surprise and amuse .

Bayesian Approach

To calculate the likelihood of a review being positive based on a single word:

- » w is a single word in the review
- » $P(\text{positive} | w) = P(w | \text{positive}) * P(\text{positive}) / P(w)$

$P(\text{positive})$ is fraction of reviews that are positive in the collection of reviews (prior likelihood of a review being positive)

$P(w)$ is the prior likelihood of seeing w

- » $P(w) = \text{freq. of } w \text{ in all reviews} / \text{total number of words in all reviews}$

$P(w | \text{positive})$ the chance of seeing w in a positive review

$P(w | \text{positive}) = \text{count of } w \text{ in positive reviews} / \text{total words in positive reviews}$

Naive Bayes Algorithm

- This is the simplest machine learning Algorithm

- Assumes that all features (words) are independent.

$$P(Pos|w1, w2, \dots, wn) = P(w1, w2 \dots wn | Pos) * P(Pos) / P(w1, w2 \dots wn)$$

$$P(w1, w2 \dots wn | Pos) = P(w1 | Pos) * P(w2 | Pos) \dots * P(wn | Pos)$$

$$P(w1, w2 \dots wn) = P(w1) * P(w2) * \dots * P(wn)$$

To simplify, calculate:

$$Score(Pos) = P(Pos) * P(w1 | Pos) * P(w2 | Pos) * \dots * P(wn | Pos)$$

$$Score(Neg) = P(Neg) * P(w1 | Neg) * P(w2 | Neg) * \dots * P(wn | Neg)$$

Return the sentiment with higher score

Naive Bayes Probability Distribution

$$\text{Score}(\text{Pos}) = P(\text{Pos}) * P(w1|\text{Pos}) * P(w2|\text{Pos}) * \dots * P(w_n|\text{Pos})$$

$$\text{Score}(\text{Neg}) = P(\text{Neg}) * P(w1|\text{Neg}) * P(w2|\text{Neg}) * \dots * P(w_n|\text{Neg})$$

If you want a probability distribution (to determine confidence of classification, for example)

$$P(\text{Pos} | w1..w_n) = \text{Score}(\text{Pos}, w1..w_n) / (\text{Score}(\text{Pos}, w1..w_n) + \text{Score}(\text{Neg}, w1..w_n))$$

$$P(\text{Neg} | w1..w_n) = \text{Score}(\text{Neg}, w1..w_n) / (\text{Score}(\text{Pos}, w1..w_n) + \text{Score}(\text{Neg}, w1..w_n))$$

Naive Bayes Probability Distribution

Naive Bayes can be used for any text classification task,
For example: email spam filter:

$$\text{Score}(\textit{Spam}) = P(\textit{Spam}) * P(w1|\textit{Spam}) * P(w2|\textit{Spam}) * \dots * P(w_n|\textit{Spam})$$

$$\text{Score}(\textit{Nospam}) = P(\textit{Nospam}) * P(w1|\textit{Nospam}) * \\ P(w2|\textit{Nospam}) * \dots * P(w_n|\textit{Nospam})$$

What if you have more than two classes? Positive, Negative and Neutral